


Day 2: Neural Networks

Guillem & Roderic Guigo Corominas


24-08-2021

SMTB

Neural Networks



Neural Networks




Loss function

- Mean Squared Error

$$\mathcal{L}(\theta) = MSE = \frac{1}{n} \sum_i (y_i - f_{\theta}(x_i))^2$$



- Cross Entropy

$$\mathcal{L}(\theta) = H = - \sum_i y_i \log(f_{\theta}(x_i))$$



Layer abstraction

- Different types of layers: Fully connected, convolutional, recurrent
- **Convolutional neural networks (CNN)**, recurrent neural networks (RNN), autoencoders, GANs, ...
- Layers of different types can be stacked to obtain new functionalities



Why do we need multiple layers?


- **Universal approximation theorem:** The standard multilayer feed-forward network with a single hidden layer, which contains finite number of hidden neurons, is a universal approximator among continuous functions on compact subsets of \mathbb{R}^n , under mild assumptions on the activation function
- We might need an extremely large number of neurons in the hidden layer to model our function
- Deep neural networks are easier to train

Train, test & validation sets

- Dataset with 80.000 labeled images
- Splits: 70% train/20% validation/10% test: 56.000 in training set, 16.000 in validation set & 8.000 in test set
- The **training set** is the data that we use to fit our model
- The **validation set** is used to adjust the **hyperparameters** of the model
- The **test set** is used to test the performance of our model in unseen data

Overfitting & Underfitting

- The number of parameters in our model is related to the dimensionality of our data and the number of examples in our training set
- Overfitting occurs when our model fits "too well" the training set, so it doesn't capture the true trend in our data



Bias/Variance tradeoff

- A highly biased model will likely have low variance, but it will not be able to fit the data properly
- If the bias is too low our model can overfit the training data and it will not do well on new data
- Ideally, we want low bias & low variance
- Regularization

Regularization

- Regularization techniques are used to avoid overfitting and improve generalization of the model
- A common technique is L2 regularization or **weight decay**. It prevents weights in the neural network from becoming too large by adding a penalty to the loss function


$$E = (\theta) + \lambda \sum_j^p \theta_j$$

- **Dropout** ($0 < p < 1$)


Hyperparameters

- Number of hidden layers
- Number of nodes in each layer
- Loss function
- Weight initialization
- Weight decay (λ)
- Dropout
- Learning rate (α)
- Activation functions
- Batch size
- Number of epochs


Learning rate




Activation functions




(a)



(b)



(c)

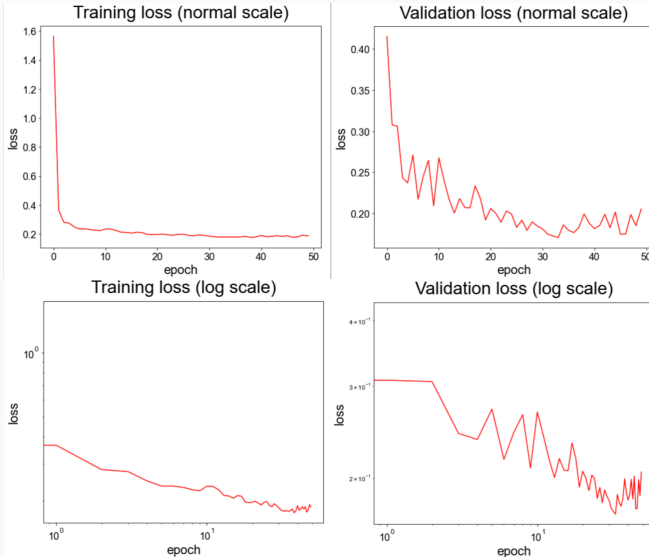


(d)

Batch size

- We (usually) don't train with the entire dataset at once, nor with a single example at a time
- Instead, we select a subset of examples in our training set and we compute the gradient after we calculate the error of all of the examples in the subset. This subset is called a **batch**
- Batch size: 1, 16, 64, 256...

Epochs



Hyperparameters

- Number of hidden layers
- Number of nodes in each layer
- Loss function
- Weight initialization
- Weight decay (λ)
- Dropout
- Learning rate (α)
- Activation functions
- Batch size
- Number of epochs

Train a Neural Network